

# Coordination and Agreement: Distributed Mutual Exclusion

Main Reference: Ch. 15 of Coulouris et.al

Supplementary: Ch. 12 of Tanenbaum & Steen



# Aim:


- Understand the problems of coordination and agreement in distributed systems; i.e. how processes coordinate their actions and agree on shared values.
- Study algorithmic techniques for distributed mutual exclusion and election problems, consensus and related problems.

# Outline:

- **Distributed mutual exclusion:** the central server algorithm, ring-based algorithm, algorithm using multicast and logical clocks, Maekawa's voting algorithm.
- **Elections:** ring-based election algorithm, bully algorithm.
- **Consensus:** Byzantine agreement, consensus using failure detectors, consensus using randomization.

# Agreement in Pepperland

- Apple and Orange (divisions of Pepperland Army) vs. Enemy
- Need to agree on: who will lead the charge and when to charge
- Communicate via messengers
- Asynchronous vs Synchronous Pepperland



# Failure detection in Pepperland

- As long as a division is not yet defeated, it will keep sending messengers to the other division.
- In async system, hard to determine if a division has defeated: a messenger may take ages to reach the other division.
- In sync system, absence of messenger indicates defeat of the other division.

# Impossibility of reaching agreement in the presence of failures

- A messenger might get caught by the enemy on it's way
- Message will not get across (either to charge or surrender), hence unable to reach an agreement.
- No protocol that guarantees agreement between Pepperland divisions can exist if messengers get captured.

# Assumptions


- Asynchronous DS (no timing assumptions) vs synchronous DS (there is a bound on the max message transmission delay)
- Begin with algorithms that tolerate no failures BUT consider how to deal with failures.

# Failure Assumptions

- Assume that each pair of processes is connected by reliable channels.
- No process failure implies a threat to the other processes' ability to communicate.
- In a sync system, a reliable channel delivers each message within a specified time frame.
- Any failed link/router will eventually be repaired.



# Failure detectors

- Knowing when a processor has crashed.
- Two types of detectors: Reliable and unreliable.
- Unreliable: Unsuspected/Suspected states 
- Reliable: Unsuspected/failure states
- To cope with failures, we must detect it.

# Distributed mutual exclusion

- Motivation – critical sections
- Model and requirements
- Evaluation criteria
- Central server algorithm
- Ring-based algorithm
- Algorithm using multicast & logical clocks
- Maekawa's voting algorithm
- Consideration for fault tolerance

# Motivation-Critical sections



- Collection of processes share resources
- When accessing shared resources (critical section), must ensure consistency and prevent interference.
- Need for distributed mutual exclusion
- Solution must be based solely on message-passing: cannot use shared memory.

# Model and Requirements

- A system of  $N$  processes  $p_i$ ,  $i=1,2,\dots,N$
- Assumptions:
  - Asynchronous system
  - Processes do not fail
  - Message delivery is reliable
- Requirements:
  - Safety: at most one process may execute CS at a time (mutual exclusion)
  - Liveness: requests to enter/exit CS eventually succeed (no deadlock, no starvation)
  - Ordering: if one request to enter CS happened before another, entry to CS is granted in that order.


# Evaluation Criteria

- **Consumed bandwidth:** proportional to number of messages sent in each entry-CS and exit-CS operations
- **Client delay:** time spent at entry-CS and exit-CS operations (worst-case)
- **Synchronization delay:** time between one process exiting CS and another process entering CS

# Central server algorithm

- A server grants permission to enter CS (via token)
- To enter CS: send request to server and wait until it replies with a token.
- To exit CS: send token to server
- Server grants token if no process holds it, else queue the request. A FCFS queue of requests is maintained by server.

# Server managing a mutual exclusion token for a set of processes



# Evaluation of algorithm


- Enter CS takes two messages ( a request followed by a grant)
- Exit CS takes 1 release message
- Synchronization delay: time taken for a round trip; I.e. a release message to the server followed by a grant message to the next process to enter CS.



# Ring-based algorithm

- Processes arranged in a logical ring
- Each process has a communication channel to the next process in the ring.
- Token is passed from process to process in a single direction
- If process that receives token does not want to enter CS, it passes token to the next process. Otherwise, it retains token until exiting CS.

# A ring of processes transferring a mutual exclusion token



# Evaluation of algorithm

- Bandwidth: Continuously consumed, except when a process is in CS.
- Requesting process delay between 0 messages (when it has just received the token) and N messages (when it has just passed on the token)
- To exit CS requires 1 message.
- Synchronization delay between exit CS and next enter CS is 1 to N message transmissions.

# Algorithm using multicast and logical clocks

- Due to Ricart and Agrawala(1981)
- Multicast a CS request to all  $N-1$  peers, request is granted when all peers reply.
- Each process keeps a Lamport clock; a monotonically increasing software counter whose value has no relationship with any physical clock. Each process keeps it's own logical clock which it uses to apply Lamport timestamp of events.

# Algorithm using multicast and logical clocks-contd

- Request message is of the form  $\langle T, p_i \rangle$ , where  $T$  is the sender's timestamp and  $p_i$  is the sender's identifier.
- Process states:

RELEASED(outside CS)

WANTED(wanting entry to CS)

HELD(in CS)

# Ricart and Agrawala's algorithm

- If process requests entry and state of all other processes is RELEASED, all the process will reply and entry is granted.

*On initialization*

*state* := RELEASED;

*To enter the section*

*state* := WANTED;

Multicast *request* to all processes;

*T* := request's timestamp;

*Wait until* (number of replies received =  $(N - 1)$ );

*state* := HELD;

request processing deferred here

# Ricart and Agrawala's algorithm

- If some processes is in state HELD, that process will not reply until it exits CS.


*On receipt of a request  $\langle T_i, p_i \rangle$  at  $p_j$  ( $i \neq j$ )  
if ( $state = \text{HELD}$  or ( $state = \text{WANTED}$  and  $(T, p_j) < (T_i, p_i)$ ))  
then  
    queue request from  $p_i$  without replying;  
else  
    reply immediately to  $p_i$ ;  
end if  
To exit the critical section  
 $state := \text{RELEASED}$ ;  
reply to any queued requests;*

# Ricart and Agrawala's algorithm

- If two or more request entry at the same time, whichever process request bears the lowest timestamp will be the first to collect  $N-1$  replies, granting it next entry.
- If process bear equal timestamps, requests are ordered according to the processes identifier.



# Multicast synchronization





# Evaluation of algorithm

- To enter CS takes  $2(N-1)$  messages;  $N-1$  to multicast request, followed by  $N-1$  replies.
- Synchronization delay: 1 message transmission.

# Maekawa's voting algorithm

- In order for a process to enter CS, not necessary for all peers to grant it access.
- Sufficient to obtain permission from subsets of peers, as long as the subsets used by any two processes overlap.
- Processes vote for one another to enter CS. A candidate must collect sufficient votes to enter.

# Maekawa's algorithm


- Each process is associated with a voting set  $V_i$
- To enter CS, process  $p_i$  sends request to  $K-1$   members of  $V_i$  where  $K=|V_i|$
- $P_i$  cannot enter CS until it has received  $K-1$   reply messages.
- When a process  $p_j$  in  $V_i$  receives  $p_i$ 's request, it replies immediately, provided it's not in HELD state or has replied (voted) since it last received a release message.

# Maekawa's algorithm - contd



- Else, it queues request in order of arrival, but does not reply.
- When it receives a release message, it removes the head of the queue and votes.
- To leave CS,  $p_i$  sends a release message to all  $K-1$  members of  $V_i$

# Maekawa's algorithm-deadlock prone

Each process has received 1 out of 2 votes, so cannot proceed.



# Consideration for Fault Tolerance

- Lost messages: None of the algorithms can tolerate loss of messages.
- process crash:
- Ring-based algorithm-NO WAY 
- Maekawa's algorithm – OK, if not in voting set 
- Central server, OK if neither holds or requested token. 