

Bachelor of Computer Science (Hons) Year-2 SEP2023

Welcome to Intelligent Systems

CAI3204N

Learning Objectives

- ❑ At the end of the course, students will be able to:
 - ❑ CO1: Identify the types of problem that are amenable to "intelligent" solutions.
 - ❑ CO2: Compare and contrast the various intelligent system techniques to solve such problems.
 - ❑ CO3: Select and apply appropriate intelligent techniques to a given problem.
 - ❑ CO4: Critically discuss intelligent system research issues and their applications.

Artificial Neural Networks

Revisit : More Examples

Learning Objectives

- ❑ At the end of the course, students will be able to:
 - ❑ CO1: Identify the types of problem that are amenable to "intelligent" solutions.
 - ❑ CO2: Compare and contrast the various intelligent system techniques to solve such problems.
 - ❑ CO3: Select and apply appropriate intelligent techniques to a given problem.
 - ❑ CO4: Critically discuss intelligent system research issues and their applications.

Neural Networks Problems perspective

Today's Overview

- Connectionism
- Biological Brains
- Artificial Neurons
- The Perceptron
- Learning
- Limits of Learning
- Artificial Neural Networks

http://en.wikipedia.org/wiki/Artificial_neural_network

Why ANNs

- Artificial Neural Networks are a machine learning technique inspired by the mechanism underpinning human learning
 - Lots of different types
 - Lots of different implementations
 - Lots of different applications
- If they work like humans, maybe they can learn to behave like humans?

Connectionism

Cognitive : (Complex Connection between nodes)

- An approach to Artificial Intelligence that seems to challenge the **Physical Symbol System Hypothesis**
- Provokes lots of interesting philosophical questions
- That we shall ignore
 - <http://plato.stanford.edu/entries/connectionism>

Brains

- Brains are the dominant organ of the **central nervous system** that **controls all intelligent behaviour**.
- Also much **unintelligent behaviour**: breathing, heart rate, blood pressure, body temperature, *etc*
- Connected to all senses and muscles throughout the **body via spinal cord**


Brain Activity

- Human adult brain weighs ~3lb. *ie* about 2% of total body weight
- But consumes 25% of total energy (60% in babies!)
- Which is why you should wear a hat when its cold
- *What is going on in there?*

Neurons



- Brains made up of neurons (and supporting glial cells)
- Neurons characterised by short and long branching connections (dendrites and axons)
- Human brain has 100 billion neurons each connected to 10,000 others

Structure of a Typical Neuron



Brain Activity




- Neurons are electrically excitable
- Gather charge from dendrites
 - (Or from sensory nerves, optic nerve, etc)
- When enough charge is gathered, neuron releases an *action potential* (spike) down the axon
 - ie when charge exceeds a threshold
- Which is transmitted on to other neurons
 - (Or on to muscle fibers)



Synapses


- The 'intelligence' in the brain lies in the connections
 - Which neurons are connected to which others
 - The strength of the connection (the *weight*)
- Connections can be *excitatory* or *inhibitory*
(*Serotonin Balance: stable mood*)
- Learning takes place by changing the strength and nature of the connections.
- A lot of psychoactive drugs work by temporarily changing the electro-chemical properties of these connections.

Hebbian Learning (Dog Training Example)




Artificial Neurons

- A cartoon model of a real neuron
- Output is a simple function of the inputs
- Exact value determined by weight of each connection, and an (optional) threshold value
- Present inputs in turn and find outputs
- Adjust weights to get behaviour we want



Activation

- Each input receives a value
- The inputs are multiplied by respective weights and added together
- Output (*Activation*) is a simple function of weighted input



$$\begin{aligned} sum &= \sum I_i w_i \\ &= I_1 w_1 + I_2 w_2 + I_3 w_3 \end{aligned}$$

$$out = f(sum)$$

Inputs to Neural Networks

- Inputs
 - Must be numeric
 - In the range $[-1,1]$ (usually)
 - Often requires *normalisation*

Name	Values	Normalisation	Normalised Input
Rainfall	0, 5, 50mm	Divide by 100	0, 0.05, 0.5
Temperature	-5, 0, 20C	Divide by 100	-0.05, 0, .2
Stock Type	Tech, Oil, Telecoms	1 binary input for each class	(1,0,0), (0,1,0), (0,0,1)

Activation Functions


- Different types of activation function produce different types of output.

Let's do this way!

- Can picture activation function as a graph of output against input sum
- The type of activation function depends on what kind of output we want
 - Binary (1 or 0, *yes* or *no*, *A* or *B*)
 - Continuous (any number)
 - Continuous range (any number 0-1)


Binary Step Function

- Produces **binary** output (0 or 1)
- Good for **yes/no** type questions
- *Output = 1 if $sum > 0$, 0 otherwise*




Bipolar Step Function

- Or *bipolar* output: -1 or 1
- $output = +1$ if $sum > 0$, -1 otherwise
- (Very similar to binary, but occasionally learns better)



Identity Function


- Produces continuous output in a wide range
- $output = sum$



Sigmoidal (Squashing) Function


- Produces continuous output in a limited range (0,1)
- Two *asymptotes*:
 - $output \rightarrow 1$ as $sum \rightarrow \infty$
 - $output \rightarrow 0$ as $sum \rightarrow -\infty$

$$output = \frac{1}{1 + e^{-sum}}$$




Thresholded Step Function

- Often useful to add a threshold
- $output = 1$ if $sum > \theta$, 0 otherwise, or
- $output = 1$ if $sum - \theta > 0$, 0 otherwise




Example Networks

Company			<i>Buy/Sell?</i>
Age	P/E	Type	
12	20	Tech	<i>Buy</i>
3	15	Oil	<i>Sell</i>
5	30	News	<i>Buy</i>
20	10	News	<i>Sell</i>
10	50	Telecoms	<i>Sell</i>



Previous Weather			Future Weather		
Temp	Rain	Wind	Temp	Rain	Sun
10C	5mm	20mph	12C	0mm	3hrs
12C	0mm	10mph	15C	2mm	0hrs
2C	15mm	45mph	5C	4mm	12hrs
-3C	2mm	5mph	0C	15mm	5hrs
8C	1mm	15mph	-4C	4mm	6hrs




f_1, f_2 = linear function
 f_3 = sigmoidal function

The Perceptron

- The simplest form of neural network
- Developed by Rosenblatt, 1957
- Used as a *decision system*
 - ie a two-class classifier
- Thresholded step activation function
- Binary (or sometimes bipolar) output
- <http://en.wikipedia.org/wiki/Perceptron>

Perceptron Thresholds

- Threshold can be represented as a weighted input




$$\begin{aligned} sum &= I_1W_1 + I_2W_2 + I_3W_3 + \dots \\ &= \sum I_iw_i \end{aligned}$$

$$\begin{aligned} out &= 1 \text{ if } sum - \theta > 0 \\ &0 \text{ otherwise} \end{aligned}$$

Perceptron Thresholds


- Threshold can be represented as a weighted input




$$\begin{aligned} sum &= I_1W_1 + I_2W_2 + I_3W_3 + \dots - \theta \\ &= \sum I_iw_i - \theta \end{aligned}$$

$$\begin{aligned} out &= 1 \text{ if } sum > 0 \\ &0 \text{ otherwise} \end{aligned}$$

Perceptron Example



Perceptron Example



I_1	I_2	I_3	$sum-\theta$	out
0.0	1.0	0.5	0.1	1
-0.2	0.0	-1.0	-1.3	0
0.2	-0.5	0.2	0.2	1
-0.8	0.2	-0.6	-1.44	0

Perceptron Learning

- How do we get a perceptron to learn to classify examples correctly?
- Present each of the examples in the training set
- See what output you get
- Adjust the weights to get the output 'more right'
- Until it does what you want

Training a Perceptron

1. Start weights **at random**
2. Present **inputs and calculate** outputs
3. Find error compared with desired output
4. Adjust weights
5. Repeat 2-4 until:
 - *Either* got the outputs you want
 - *Or* results not getting any better
6. Then use network to make predictions

Learning Rule

- How to adjust the weights?
 - If $\text{input} > 0$:
 - If the answer is too big, reduce the weight
 - If too small, increase it
 - *but if input < 0*:
 - If the answer is too big, *increase* the weight
 - If too small, *reduce* it
 - Only increase/decrease by a little bit at a time

Learning Rule


error = output – target 

rate = 0.2 (for example) 

$$W_{new} = W_{old} - (error \times input \times rate)$$




Learning Example - Error



$$error = output - target$$

I_1	I_2	I_3	$Sum-\theta$	out	<i>target</i>	<i>error</i>
0.0	1.0	0.5	0.1	1	1	0
0.2	-0.5	0.2	0.2	1	0	1
-0.2	0.0	-1.0	-1.3	0	1	-1
-0.8	0.2	-0.6	-1.44	0	0	0

Learning Example



$$W_{new} = W_{old} - (error \times input \times rate)$$

rate=0.2

I_1	I_2	I_3	Sum	out	target	error
0.2	-0.5	0.2	0.2	1	0	1

error = 1 ie output is too large

- $I_1=0.2$, so decrease W_1
- $I_2=-0.5$, so increase W_2
- $I_3=0.2$, so decrease W_3
- $I_\theta=-1$, so increase θ

$$W_{1new} = 0.5 - (1 \times 0.2 \times 0.2) = 0.46$$

$$W_{2new} = -0.2 - (1 \times -0.5 \times 0.2) = -0.1$$

$$W_{3new} = 1 - (1 \times 0.2 \times 0.2) = 0.96$$

$$\theta_{new} = 0.2 - (1 \times -1 \times 0.2) = 0.4$$

I_1	I_2	I_3	Sum	out	target	error
0.2	-0.5	0.2	-0.066	0	0	0



Training Regimes


- Suppose there are three examples in training set
 - {A,B,C}
 - Present each example repeatedly until successful on each
 - Or Present complete set in turn until successful on all?
- Which training regime would be best?
- The latter, since by the time we have trained on B and C, may have 'forgotten' how to do A

A	x		A	x
A	x		B	x
A	x		C	✓
A	✓		A	✓
B	x		B	x
B	x		C	✓
B	✓	OR	A	x
C	x		B	x
C	x		C	✓
C	x		...	
C	x		A	✓
C	✓		B	✓
			C	✓

Limits of the Perceptron


- Perceptron was invented in 50s
- Seemed to promise much
- But in 1969 **Minsky and Papert** proved there were limits on what it could learn...

Straight Line in the Plane




But this form of the equation cannot describe a vertical line!

Or, more generally



Linear Separation




Linear Perceptron

- Equation for the line:

- $ax + by - c = 0$
- Two regions
 - $ax + by - c > 0$
 - $ax + by - c < 0$


- Equation for the perceptron:

- $w_1in_1 + w_2in_2 - \vartheta = \text{sum}$
- Two possible outputs
 - $w_1in_1 + w_2in_2 - \vartheta > 0$
 - $w_1in_1 + w_2in_2 - \vartheta < 0$



*ie They are of the same form!
(And this applies to >2 inputs/dimensions)*

Linear Perceptron




- A single perceptron can learn to solve a problem if it is *linearly separable*
- *ie* if the two classes can be separated by a single straight line (or hyperplane)

Example Problems

- Logical **AND**
- Linearly separable?


x	y	AND
0	0	0
0	1	0
1	0	0
1	1	1



Example Problem

- Logical **OR**
- Linearly separable?

x	y	OR
0	0	0
0	1	1
1	0	1
1	1	1




Example Problem

- Logical XOR
- Linearly separable?

x	y	XOR
0	0	0
0	1	1
1	0	1
1	1	0

Logical XOR



XOR (and many other problems like it) are not linearly separable.


This scuppered research into neural networks for 20 years

Artificial Neural Networks

- Solution: combine neurons into a network
- By combining many neurons into a network can solve a wider range of more complex problems
 - Feedforward, multilayer networks
 - Recurrent (feedback) networks
 - Dynamic neural networks
- More powerful (in general), but harder to train


Multilayer Feedforward Networks

- Comprised of *layers* of simple neurons, each layer feeds into the next
- Can calculate any arbitrary function
- Well understood and can be trained using *backpropagation*
- Now applied to just about any machine learning problem




Multilayer Perceptron Example

- This 2-layer perceptron can solve the non-linearly separable XOR problem.




Backpropagation

- Invented in 1986, and made neural nets popular again.
- Problem:
 - Can use *error* to adjust input weights w_1 and w_2 for unit A
 - But what error do we use to adjust w_{3-6} for the hidden layers?
- Solution:
 - 'backpropagate' the error to update the hidden weights
- en.wikipedia.org/wiki/Backpropagation




Elman Net




- Very simple *recursive* (feedback) network
- Previous state of neuron is fed back in at next time step
- Allows for simple memory or state
- Can **distinguish** order in **sequences** of inputs
- *Eg* simple natural language processing
 - 'dog is dead'
 - 'is dog dead'
- Or **predicting time series**
 - Stock markets
 - Weather
 - Sales

Dynamic (Real Time) Neural Network



- Continuous real inputs, activation equations and outputs
- Can be used for real time control problems
 - Robotics
 - Chemical plants
 - Jet engines

Recurrent Dynamic Neural Nets



- A lot more like the real brains that control our intelligent behaviour
- Continuous, real-time inputs and outputs (eg robot motors and sensors)
- Very complex dynamics – don't know how to make them do what we want
- Cannot reliably train them – so use evolution

Summary

- Biological Brains
- Artificial Neurons
- The Perceptron
- Learning
- Limits of Learning
- Artificial Neural Networks

Stretch Break!

Second Half (NEXT WEEK)

- Human Learning
- Machine Learning
 - Supervised Learning
 - Examples
- Data Mining
 - Examples

Thank you