

COS3043

System Fundamentals

Lecture 2 (Part 1)

List of Discussion

Part 1

- OS Structure Overview
- The SPIN Approach

Part 2

- The Exokernel Approach
- The L3 Microkernel Approach

OS Structure Overview

What is Operating System Structure?

- In a nutshell, OS structure is the way the operating system software is organized with respect to the applications it serves and the underlying hardware it manages.



Question?

- Why OS structure is important?
 - Protection?
 - Performance?
 - Flexibility?
 - Scalability?
 - Agility?
 - Responsiveness?

Characteristics of Good OS Structure

Goal 1

Protection
Within and
across **users**
+ **OS** itself.






Goal 2

Performance
Time taken
to perform a
service.

Goal 3

Extensibility
Flexibility =>
NOT one size
fits all.

Goal 4

Scalability 
Performance
 when
resources .

Goal 5

Agility 
Adapting to
application
changes such
as need of
resources.

Goal 6

Responsiveness
Reacting to
external events.

Question?

- Do you think an OS including the commercial OS such as Windows, Mac and Linux can meet all the goals at the same time?
 - Protection?
 - Performance?
 - Flexibility?
 - Scalability?
 - Agility?
 - Responsiveness?

Foundation of OS Structure

| Type |
|-----------------------------|
| Monolithic Structure |
| DOS-Like Structure |
| Microkernel-Based Structure |

Monolithic Structure



Each app is on its own hardware address space – it's protect from one another.

OS Services and Device Drivers

Hardware

- *OS services => files system, network access, scheduling to CPU, virtual memory management etc.*
- *The code and data structure of the OS are contained in its own hardware address space – it's protected from the applications.*

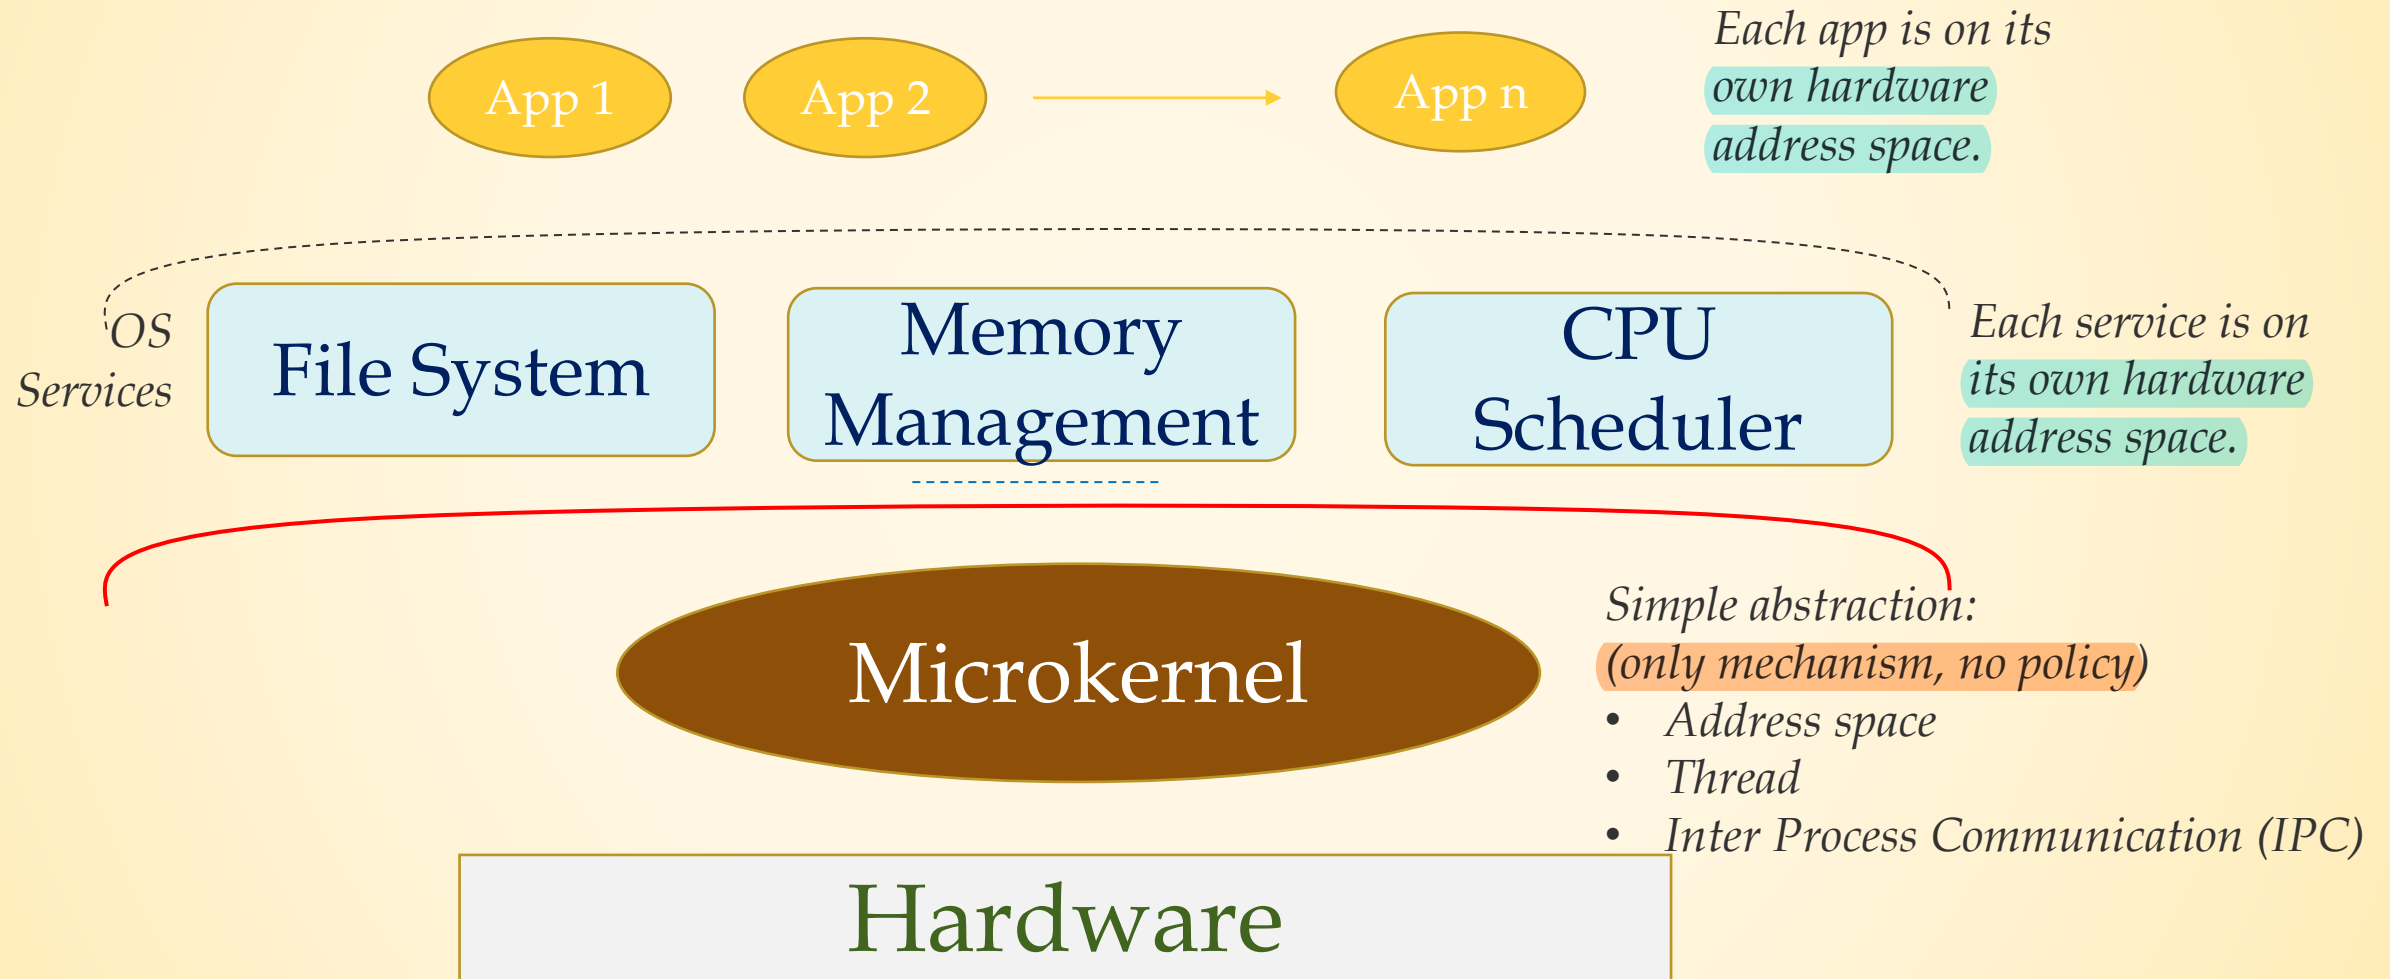
DOS-Like Structure



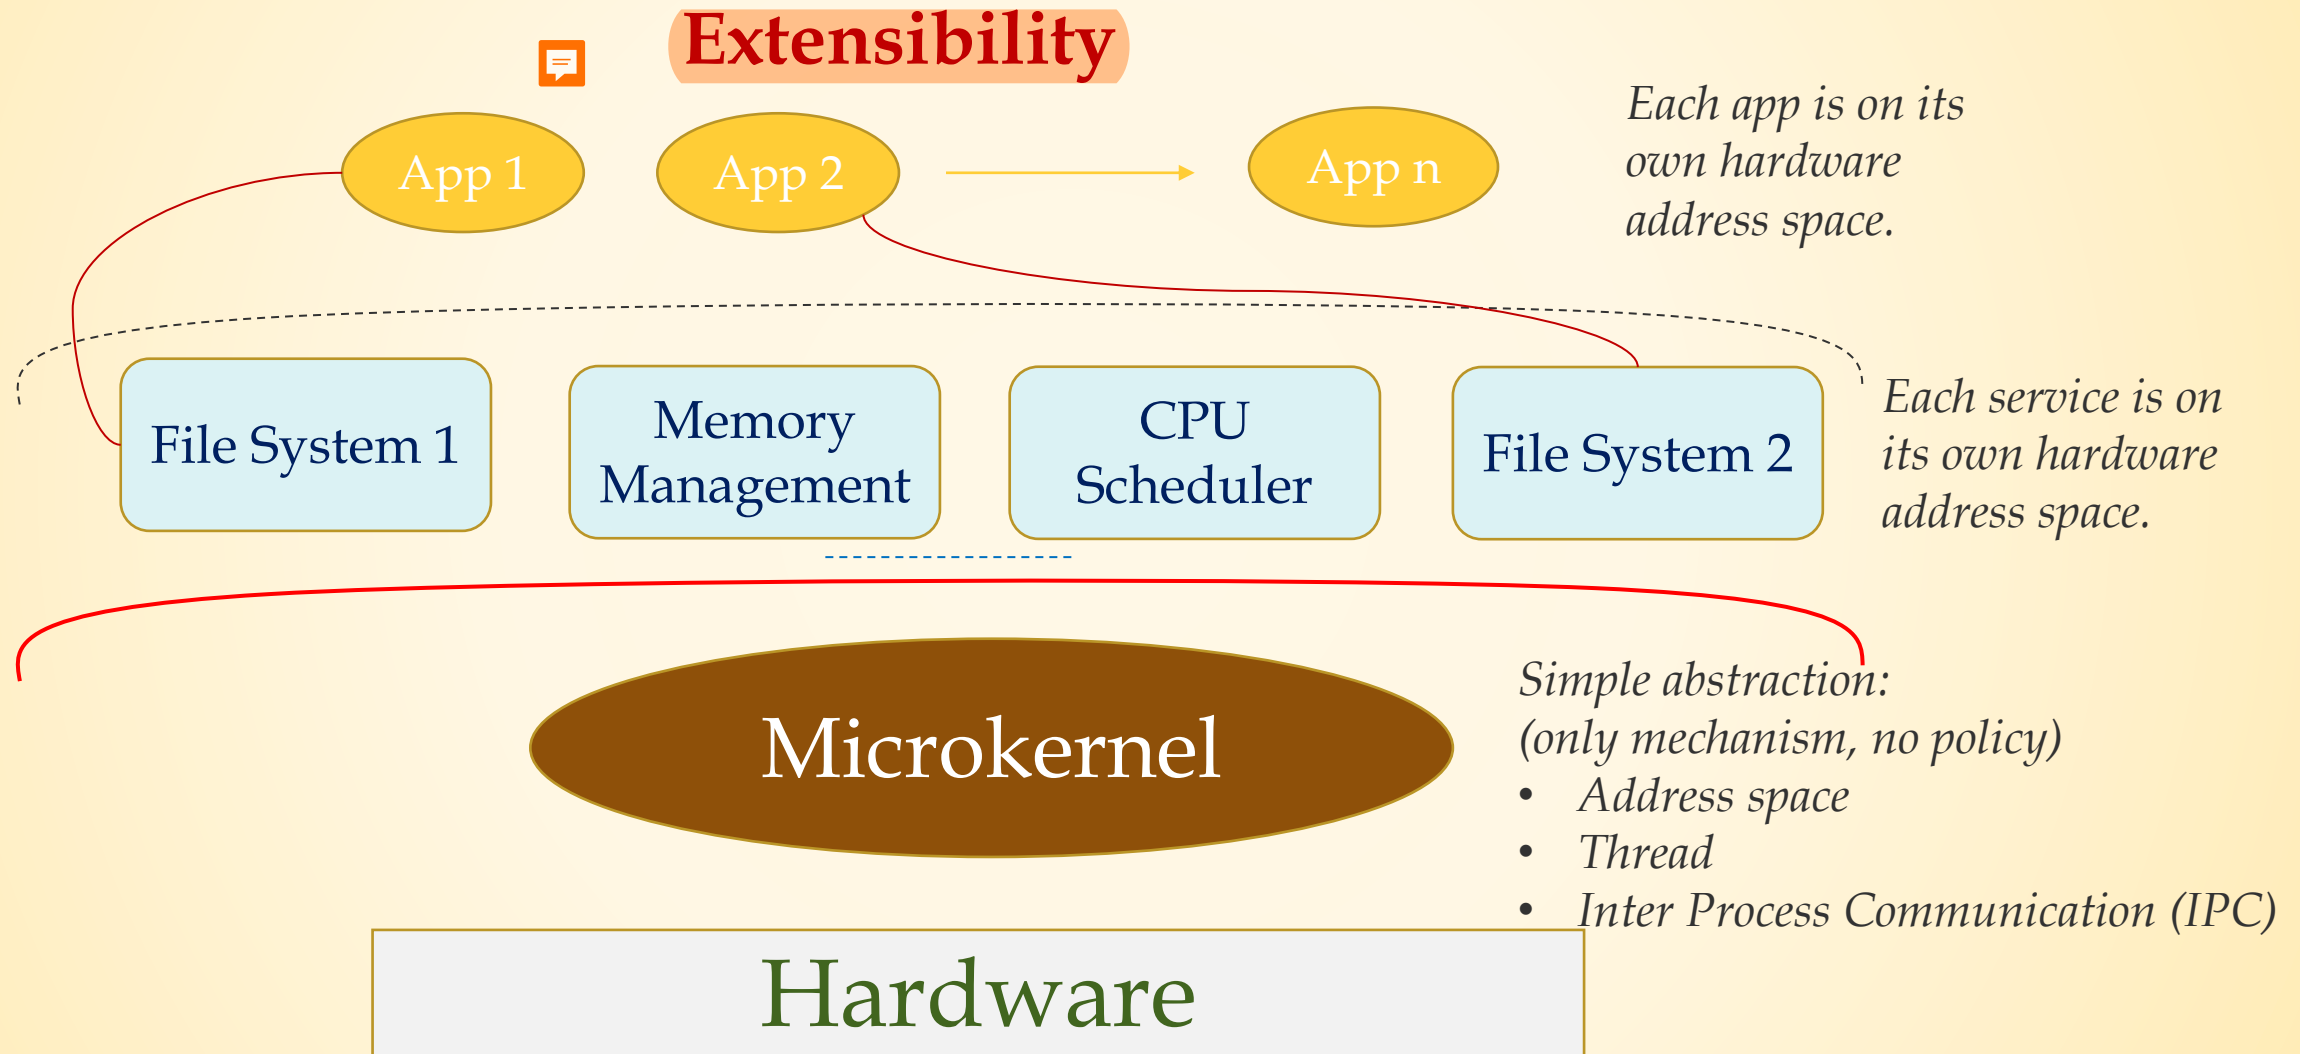
Comparison

- DOS-Like
 - Loss of protection, unacceptable for a general purpose OS.
- Monolithic
 - Loss of performance for protection?
 - Reduce performance loss by consolidation.
 - Consolidation – consolidate the components of OS in a single structure so that interaction among components can be expedited.
 - But because of the consolidation, there is no customization for different applications. But why customization?


Microkernel-Based Structure





Advantage of Microkernel-Based Structure




Shortcoming of Microkernel-Based Structure

- Potential of performance loss.
 - Border crossing
 - Change in locality 
 - User space \leftrightarrow system space copying.



Summary

- So what do we want at the end? Let's check again what do we have from the three structures...





The SPIN Approach




Introduction

- Two important premises prior to create this approach:
 - Microkernel-based design compromising on performance due to frequent border crossing.
 - Monolithic design does not lend itself to extensibility.
- So, what do we really aim for a 'perfect' OS structure?
 - Thin (like microkernel), only mechanisms, no policies.
 - Access to resources without border-crossing (like DOS).
 - Flexibility for resources management (like microkernel), without scarifying protection and performance (like monolithic).



Examples of Approach Towards Extensibility

- [Hydra OS - 1981]
 - Kernel mechanisms for resource allocation.
 - Capability based security to access resource. 
 - Resource managers as coarse-grained objects to reduce border-crossing.
 - However, because of the implementation of Capability based in managing resource, Hydra did not achieve the goal of extensibility.
- [Mach 1990s] 
 - Focused on portability + extensibility
 - Huge loss in performance because of 'portability' – painted a bad image for microkernel.

SPIN Approach Towards Extensibility

- Co-location of kernel + extensions (co-kernel for OS services)
 - Avoid border crossing.
 - But if there are same hardware address space for kernel + extensions, then does it mean it has to give in the goal of protection?
- Compiler enforced modularity
 - Use strongly typed language concept for the kernel. 
- Logical protection domains
 - Not relying on hardware address space. 
- Dynamic call binding
 - Flexibility. 

Logical Protection Domains

- Modula-3 => safety + encapsulation mechanisms 
 - Type safety, auto storage management.
 - Objects, threads, exception, **generic** interfaces.
- Fine-grained object => protection via **Capabilities** 
 - Hardware resources (ex: page frame).
 - Interfaces (ex: page allocation module)
 - Collection of interfaces (ex: the entire virtual memory sub system)
- **Capabilities** here differs from the one used for Hydra OS
 - As language supported **pointers**.

SPIN Mechanisms for Protection Domains

- Create



- Initiate with **object file contents** + export **names**.

- Resolve

- **Names** between source + target **domains**



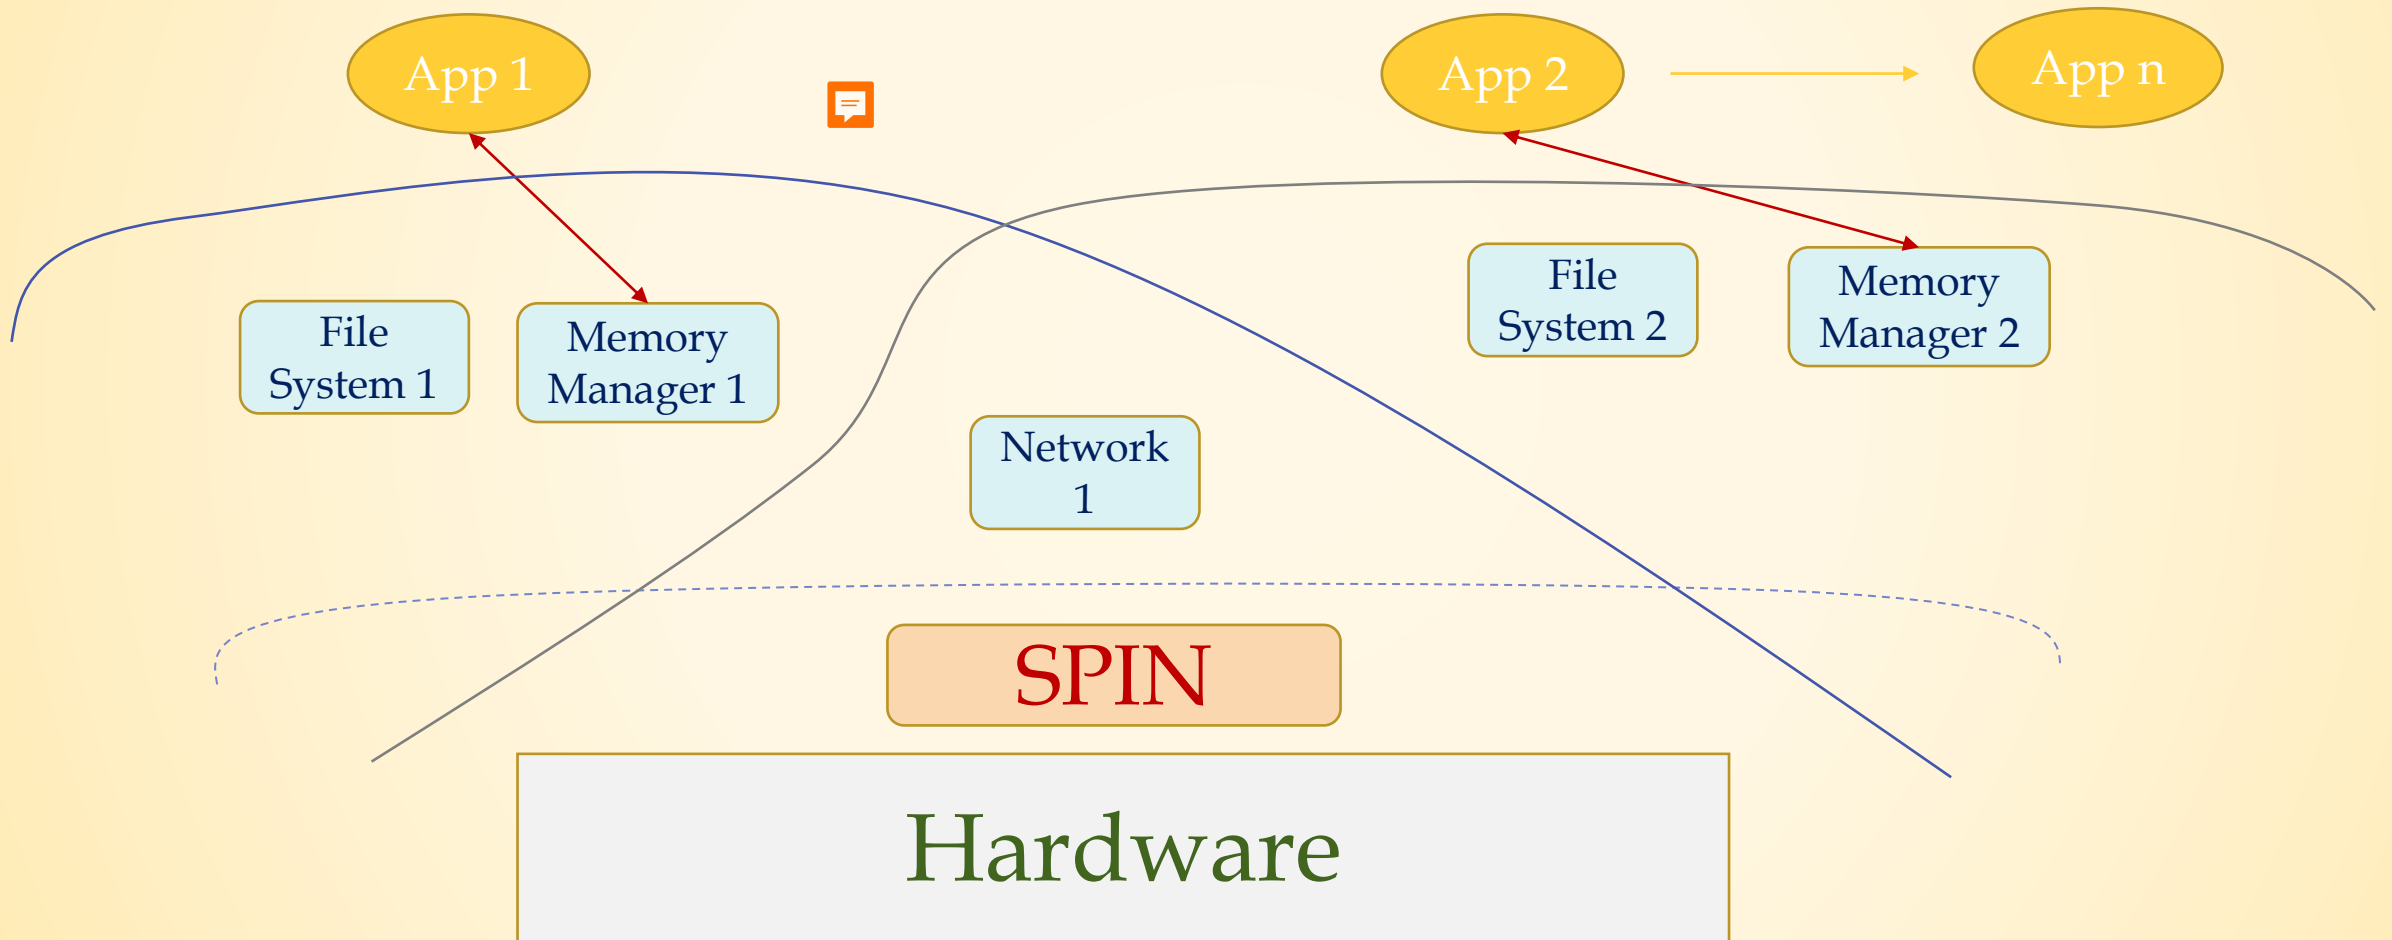
- Once resolved, resource sharing at memory speed.

- Combine

- To create an **aggregate domain**.




SPIN Approach



SPIN Mechanisms for Events

- Every OS has to answer to external events (interruption, exception)
- SPIN uses Event-Based Communication model to deal with events.



Default Core Services in SPIN

- Thus far, we know that we can build OS services from scratch as an extension in SPIN.
- However, some of the core services such as **memory management and CPU scheduling** should not be dictated by the extension.
- Thus, **SPIN provides interface procedures to** implement those **core services.**